

# SERIAL AND PARALLEL IMPLEMENTATION OF MATRIX MULTIPLICATION AND INVERSION

COURSE OF ADVANCED COMPUTER ARCHITECTURE

A.Y. 2021/2022

*Riccardo Crescenti*  
*Riccardo Doveri*

# INTRODUCTION

- The goal of this project is to implement both serial algorithms, parallelize them using open MPI and perform a performance and speed up analysis executed both in a local machine and through scalability tests on clusters built using Google Cloud Platform.
- We decided to use standard algorithm so we could operate with matrices of any dimensions. The only constraints came from mathematical point of view
- Product: the number of columns of the first matrix must be equal to the number of columns of the second matrix.
- Inversion: matrix has to be NxN, otherwise it can't be inverted. LU decomposition method was used since it has less computational cost ( $\frac{8}{3} n^3$ )

# MATRIX PRODUCT

### Algorithm

---

**Algorithm 1** Matrix Multiplication

---

INPUT: Matrices  $A$  and  $B$

OUTPUT: Matrix  $C$

```

1: for  $i = 0, 1, 2, \dots, m$  do
2:   for  $j = 0, 1, 2, \dots, p$  do
3:      $sum = 0$ 
4:     for  $k = 0, 1, 2, \dots, n$  do
5:        $sum += A_{i,k} \cdot B_{k,j}$ 
6:     end for
7:      $C_{ij} = sum$ 
8:   end for
9: end for

```

## Domain decomposition

A	A00	A01	A02	A03	B	B00	B01	B02	B03
	A10	A11	A12	A13		B10	B11	B12	B13
	A20	A21	A22	A23		B20	B21	B22	B23
	A30	A31	A32	A33		B30	B31	B32	B33
C	C00	C01	C02	C03	A TUTTI	PR. 0			
	C10	C11	C12	C13		PR 1			
	C20	C21	C22	C23					
		C30	C31	C32	C33				

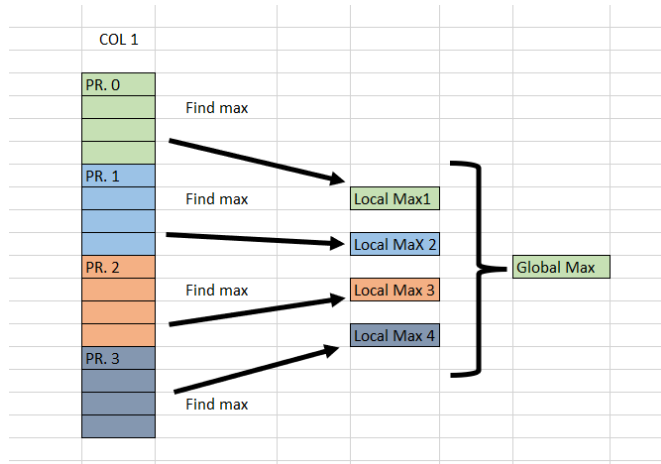
# MATRIX INVERSION

3 parts:

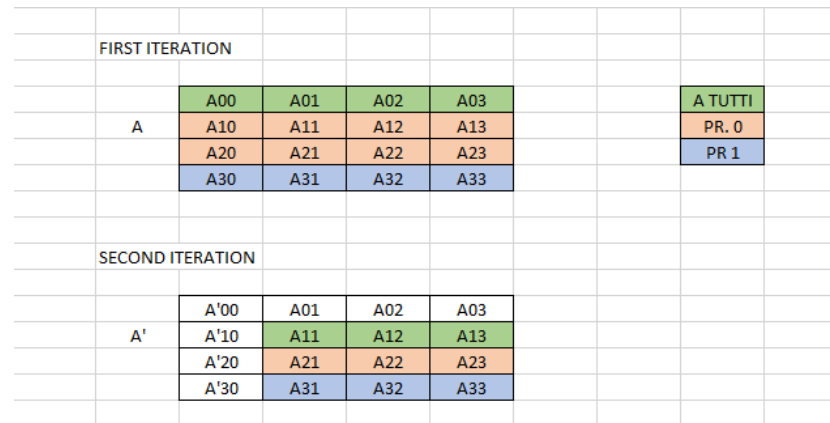
- LU decomposition with pivoting (less parallelizable, must proceed column by column)
- Check determinant
- LUP inversion (highly parallelizable, column blocks division)

# MATRIX INVERSION – domain decomposition

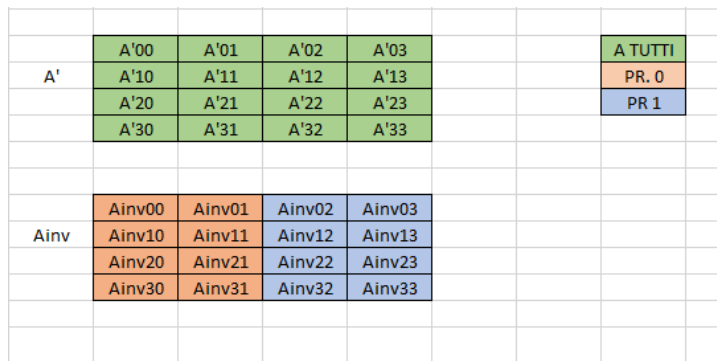
*Local maxima*



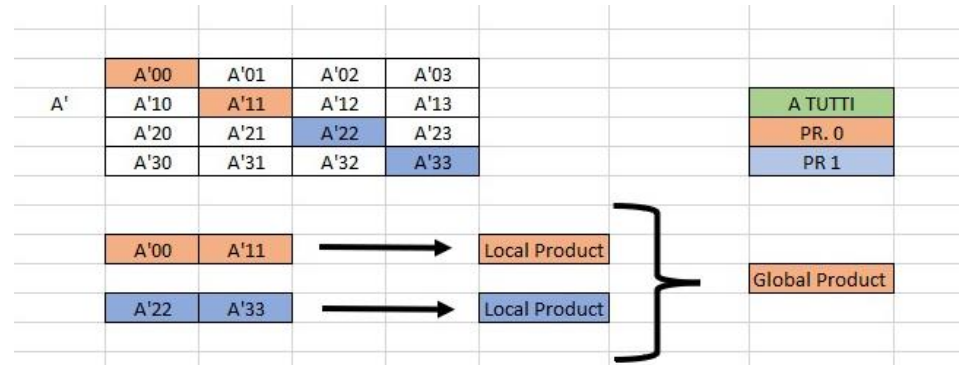
*LU decomposition*



*Inversion*



*Check determinant*



# MATRIX INVERSION – versions 2 and 3

- The inversion function from  $A'$  to  $A_{inv}$  showed improvements visible immediately
- On the other hand, the functions for calculating the maximum and determining check, for the cases tested, were not particularly efficient
- We created a v2 version of the program in which we have replaced the portion of code in parallel with the serial one.
- In the v3 version we performed the LU decomposition, maximum search, and determinant check functions in serial, keeping only the true inversion function in parallel.
- We think this is due to over-parallelization for relatively large matrices (our tests are limited to cases of  $3000 \times 3000$ ). Based on our studies in the case of even larger matrices we think that this difference will not be so deep and the v1 will overcome v3.

# ESTIMATING THE SPEED UP

- Thanks to the Amdahl's law, we can approximate the speed up for the 2 algorithms
- We estimated 97.8% of the product and about 84% of inversion
- P is the fraction of parallelizable code
- S is the not parallelizable code
- N is the number of cores

$$Speedup(N) = \frac{Time_{serial}}{Time_{parallel}(N)} = \frac{(S + P)T_{serial}}{S \cdot T_{serial} + \frac{P \cdot T_{serial}}{N}} = \frac{S + P}{S + \frac{P}{N}} = \frac{1}{S + \frac{P}{N}}$$

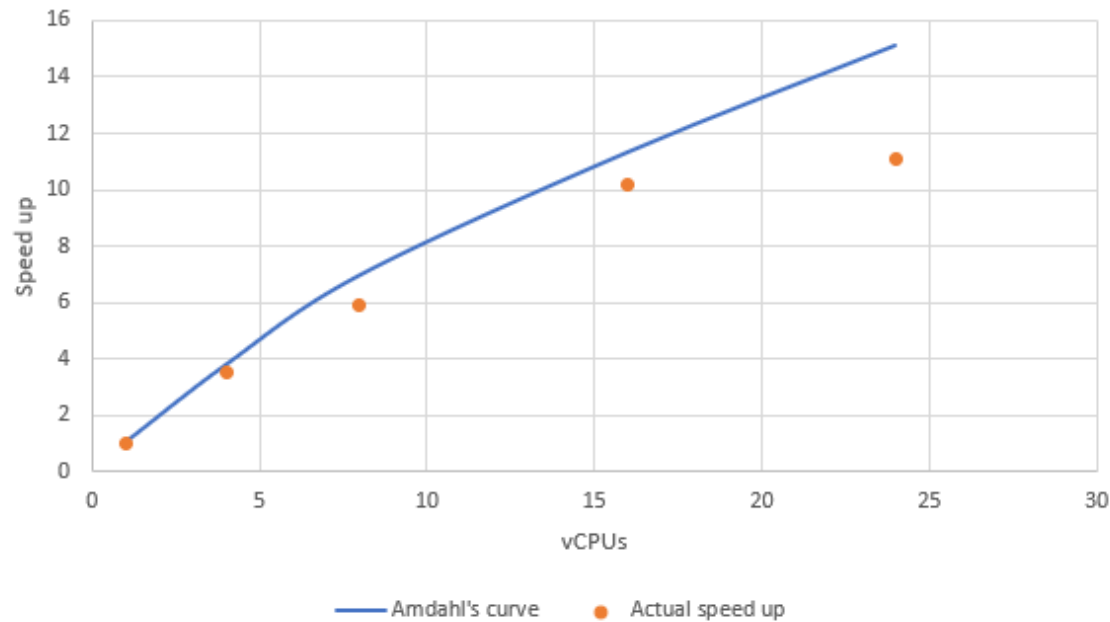
# TESTING AND DEBUGGING

- Initially on a local virtual machine (CENTOS 8), then we moved to Google Cloud Platform
- Checking the validity of the serial algorithms for small matrices with an online calculator
- Checking validity for large matrices, created with memory allocation functions
- Checking validity of parallel algorithms for a small number of cores in local
- Scalability tests on GCP
- All matrices are generated by a program we made
- Check of possible problems as file not found, incorrect matrices, incorrect dimensions or matrix degenerated

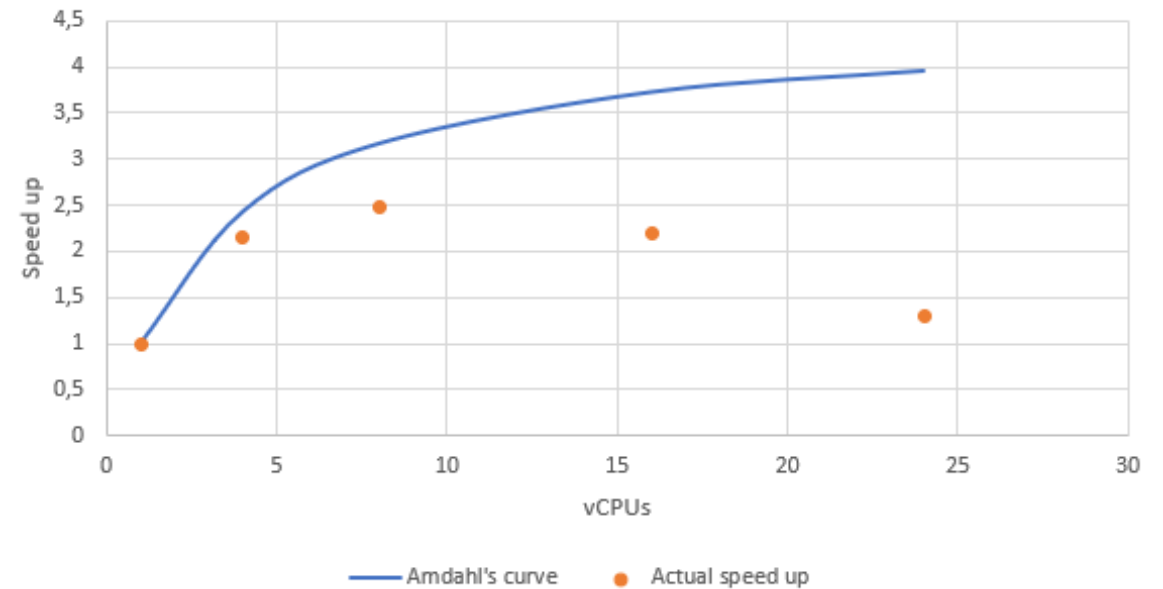


# PERFORMANCE ANALYSIS – speed up

comparison between ideal and actual speed up product

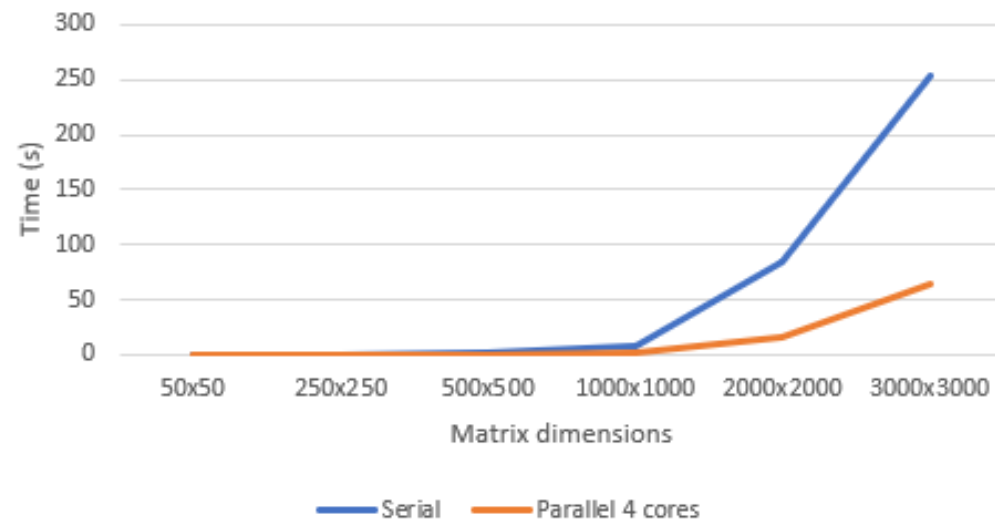


comparison between ideal and actual speed up inversion  
version 3




# PERFORMANCE ANALYSIS – local execution

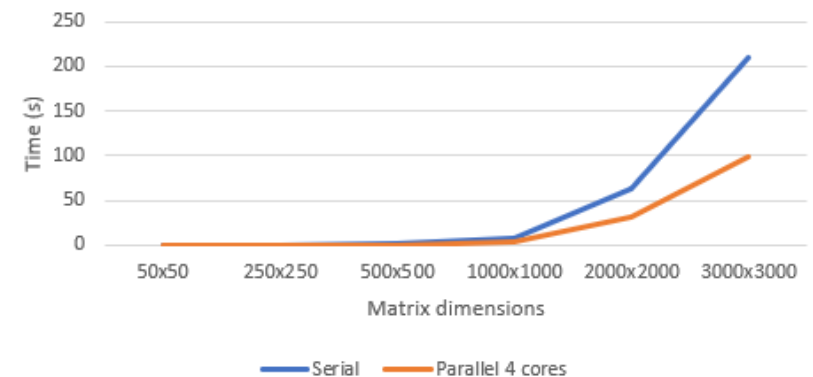
Comparison between serial and parallel product



Comparison between serial and parallel inversion  
version 1




Comparison between serial and parallel inversion  
version 3

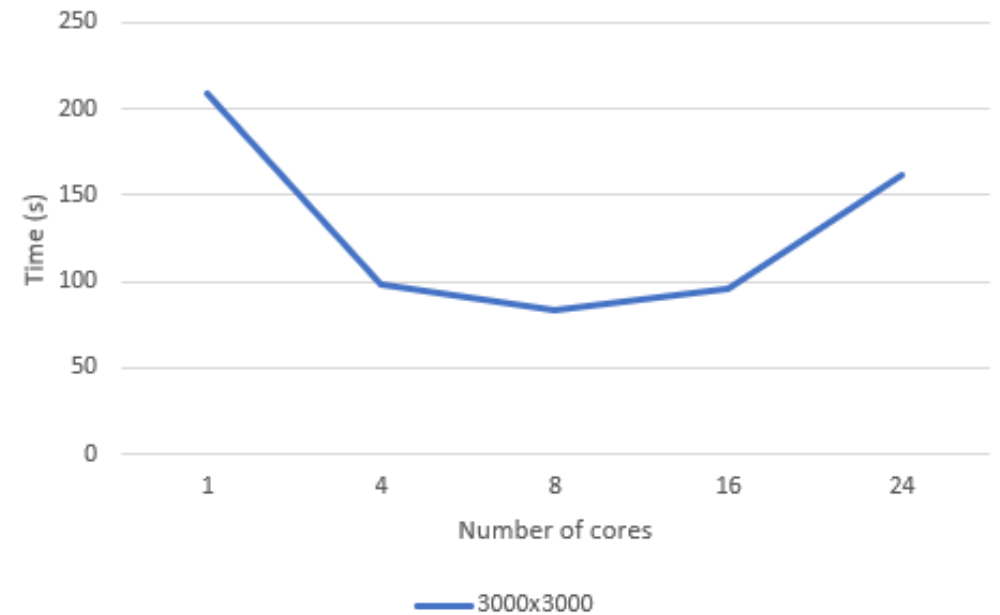


# PERFORMANCE ANALYSIS – GCP execution

Scalability test on parallel product



Scalability test on parallel inversion  
version 3

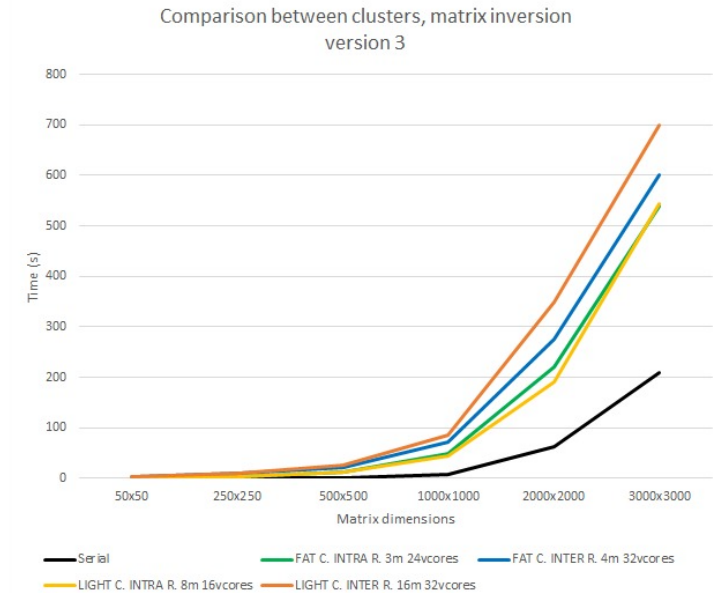

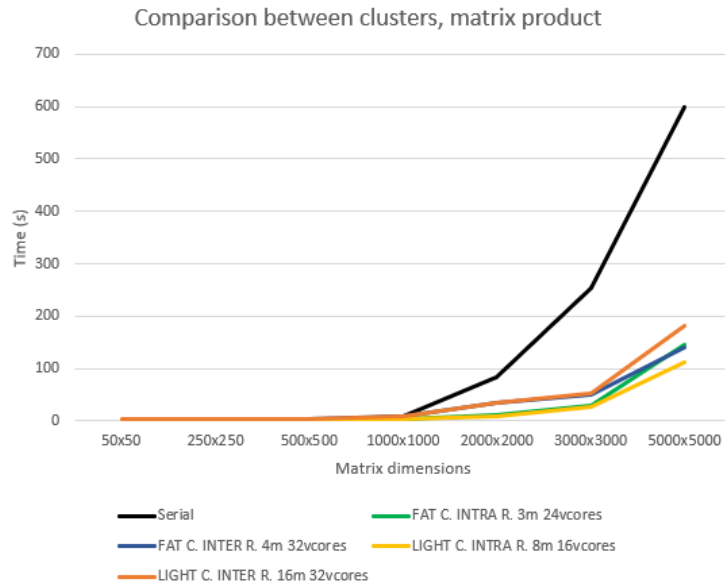


# PERFORMANCE ANALYSIS – clusters

- Fat cluster:
  - 3 machines with 8 core each, means 24 total infra-regional (US central)
  - 4 machines with 8 cores each, equal to 32 inter-regional
- Light cluster:
  - 8 machines with two cores each means 16 total infra-regional cores (US central)
  - 16 machines with two cores each equals to 32 inter-regional cores



# PERFORMANCE ANALYSIS – clusters



Second generation CPUs were selected, whose CPU platform is selected based on availability. The 8-core machines for fat clusters were configured with 36GB of RAM and 50GB of disk. The 2-core machines for light clusters were configured with 8GB of RAM and 50GB of disk.

# CONCLUSIONS

- We implemented serial matrix and inversion product algorithms
- Parallelized and performed a performance and speed up analysis
- Obtained results that reflect the values studied a priori quite well
- Best and worst solution for product/inversion/clusters were found, but it would be necessary to make tests with very large matrices.